



FedAvg vs FedProx — single seed=42, $\mu=0.01$, $E=20$, $R=150$, partition=balanced_paired_7_clients
 Δ test macro-F1 = +0.0457 Δ balanced acc = +0.0572


(A) Validation macro-F1 vs round




(B) Validation loss vs round




(C) Training loss (size-weighted, log scale)



(D) Final test metrics



(E) Per-class test F1



(F) Per-class Δ (green = FedProx better)

